

# Is Summary Useful or Not?

An Extrinsic Human Evaluation of Text Summaries on Downstream Tasks

Xiao Pu

Mingqi Gao

Xiaojun Wan



北京大學  
PEKING UNIVERSITY


# Motivation



# Motivation




# Motivation



# Motivation



# Motivation



# Methodology

---

- A summary is useful if it can facilitate users to complete a task
- 2 dimensions of usefulness: time and correctness
- We design different downstream tasks to represent diverse real-world applications of summaries:

TASK	METRIC
Question answering	Answerable, EM, F1
Classification	EM, F1
Similarity assessment	MSE, Spearman's $\rho$

# Experimental Setting -- Dataset

---



## QA

randomly collect 100 pairs of source articles and reference summaries from CNN/DailyMail (Hermann et al., 2015; Nallapati et al., 2016).

We then annotate two datasets: QA-ref (QA-pairs are written according to the reference summaries) and QA-src (QA-pairs are written according to source articles).



## Classification

randomly sample 100 news articles from New York Times Annotated Corpus (Sandhaus, 2008). Each article is paired with one or multiple tags.



## Similarity Assessment

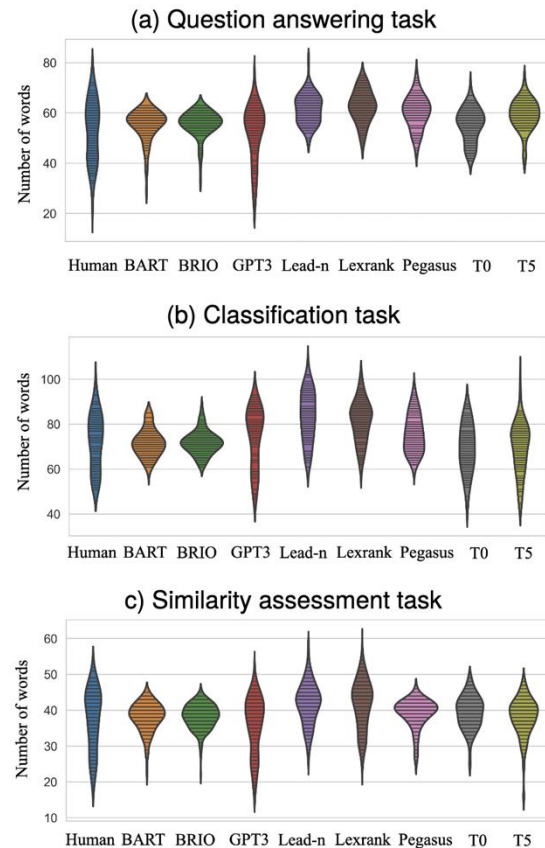
we use the SemEval-2022 Task 8 dataset (Chen et al., 2022) to collect 100 pairs of news articles with summaries and similarity scores.



# Experimental Setting – Models and Details

- Include 8 summarization models: BART, Pegasus, Lexrank, Lead-n\*, BRIO, T5, T0, GPT3
- To ensure fairness in comparing summaries across different systems, we generate summaries of similar lengths for each task
- Include 20 university students proficient in English in the experiment

\* We modify the Lead-3 setting and refer to it as the Lead-n model, which selects the first several sentences that are closest to the summary length we set.



Length of summaries from different systems in three tasks

# Experimental Setting – Platform

## A Web-based Platform for Evaluation

- Offers guidelines for annotators
- Collect experiment data, including the answers and completion time of each question
- prohibits the utilization of the copy-paste/search functionality to guarantee impartiality

The screenshot shows a web application titled "Extrinsic Eval for Text Summarization". On the left is a sidebar menu with options: HOME, QA Tasks (highlighted), Classification Tasks, and Similarity Tasks. The main content area displays a text passage about Deborah, 43, from Lanarkshire, who describes her naturally curly hair as a 'frizz nightmare'. The passage mentions her hair coloring, a salon in Glasgow, and a hair relaxing treatment. Below the passage, it says "Please answer the following questions:". There are four text input fields with corresponding questions: "Who says 'frizz nightmare' hair like Deborah's can be a battle?", "Which salon was Deborah sent to?", "Where is the salon located?", and "What treatment did Deborah have in the salon?". A green "Submit" button is at the bottom right. The top right corner shows a "Welcome" message and a user profile icon.

A screenshot of our platform

# Results: Evaluating Summaries' Usefulness

RQ1: How useful are text summaries compared to source articles?

- The use of summaries generally reduces the completion time.
- Summaries are particularly useful in classification and similarity tasks, with higher correctness on tasks.


	QA (ref-based)						QA (source-based)									
	Answerable	EM	F1	Time(seconds)			Answerable	EM	F1	Time(seconds)						
Source	0.86	0.32	0.51	280			0.89	0.51	0.7	212						
Human Summaries	0.89	+4%	0.54	+67%	0.75	+48%	94	-66%	0.54	-39%	0.27	-46%	0.4	-45%	88	-58%
All Summaries	0.52	-39%	0.22	-32%	0.33	-36%	106	-62%	0.52	-41%	0.24	-53%	0.3	-52%	83	-61%

	Classification					Similarity						
	EM		F1		Time(seconds)	MSE		Spearman's $\rho$		Time(seconds)		
Source	0.88		0.90		73	0.91		0.6		38		
Human Summaries	0.91	+3%	0.92	+2%	34	-53%	0.77	-15%	0.7	+14%	20	-47%
All Summaries	0.89	+1%	0.90	-	30	-59%	1.02	+11%	0.6	-	22	-42%

Summaries compared to source texts in the downstream tasks. The green percentages indicate that summaries are more useful compared to the source text, i.e. participants take less time or perform better. The red ones indicate less useful.

RQ2: Which summarization systems are more useful?

- Divide the summarization models into fine-tuned, zero-shot, and simple extractive.
- Fine-tuned models have higher consistency in usefulness across different tasks, and are less sensitive to differences between tasks.
- Zero-shot and simple extractive methods exhibit a varying ranking across tasks.



Average ranking of different systems on three different tasks. Each ranking is calculated by averaging the rankings over extrinsic metrics for the same task

### RQ3: What kind of summaries are more useful?

Explore how the inner features of summaries influence their usefulness in tasks

#### Inner properties and their metrics:

- Summary Style (abstractive or extractive): we employ the Ext-cvg (Extractive Fragment Coverage) to assess the extractiveness of summaries
- Grammaticality: the ratio of grammar errors in the summaries
- Sentence length: average number of words per sentence of summaries

#### Findings:

Abstractive summaries tend to be more useful for classification and similarity tasks. Grammatically correctness and shorter sentences contribute to more useful summaries in QA and similarity tasks.

	Ext-Cvg(%)	Errors(%)	Sent-Len
<b>Ref</b>	87.51	10.89	16.95
<b>BART</b>	98.83	5.13	17.91
<b>BRIO</b>	96.60	3.66	15.96
<b>GPT3</b>	93.01	2.78	24.12
<b>Lead-n</b>	100.00	9.09	28.24
<b>Lexrank</b>	100.00	10.82	29.10
<b>Pegasus</b>	98.96	5.28	17.83
<b>T0</b>	94.97	3.20	18.29
<b>T5</b>	96.44	9.17	16.18

intrinsic features of summaries from different systems


	Ext-Cvg	Errors	Sent-Len
<b>qa_EM</b>	-0.223	-0.440	-0.595
<b>qa_F1</b>	-0.291	-0.441	-0.597
<b>cls_EM</b>	-0.602	0.120	-0.090
<b>cls_F1</b>	-0.591	0.072	-0.143
<b>sim_MSE</b>	-0.641	-0.603	-0.551
<b>sim_Spearman's <math>\rho</math></b>	-0.642	-0.597	-0.507

System-level pearson correlation between intrinsic features and our extrinsic metrics

# Results: Correlation between Metrics

(1) Analyzing the relationships between our extrinsic metrics:

- Extrinsic metrics within the same task are highly correlated.
- There are only weak to moderate correlations among tasks, meaning that the tasks involved are diverse, reflecting different perspectives of usefulness



System-level Pearson correlation of our extrinsic metrics

## (2) Evaluating automatic metrics using extrinsic criteria:


- Automatic metrics can well reflect the usefulness of summaries in the QA task, but their correlations with extrinsic metrics are generally low for classification and similarity tasks.

Extrinsic Criteria Automatic Metrics	QA (ref-based)								QA (source-based)								Classification				Similarity			
	answerable				F1				answerable				F1				EM		F1		MSE		$\rho$	
	$r$	$\tau$	$r$	$\tau$	$r$	$\tau$	$r$	$\tau$	$r$	$\tau$	$r$	$\tau$	$r$	$\tau$	$r$	$\tau$	$r$	$\tau$	$r$	$\tau$	$r$	$\tau$	$r$	$\tau$
ROUGE-1	0.95**	0.71*	0.94**	0.76**	0.98**	0.86**	0.95**	0.79**	0.89**	0.64*	0.91**	0.64*	0.51	0.50	0.56	0.50	0.48	0.43	0.40	0.36	0.48	0.43	0.40	0.36
ROUGE-2	0.97**	0.79**	0.94**	0.91**	0.98**	0.93**	0.92**	0.71*	0.89**	0.71*	0.89**	0.71*	0.23	0.21	0.29	0.21	0.18	0.29	0.10	0.36	0.18	0.29	0.10	0.36
ROUGE-L	0.99**	0.93**	0.93**	0.76**	0.97**	0.79**	0.91**	0.71*	0.87**	0.71*	0.87**	0.71*	0.33	0.43	0.40	0.43	0.29	0.29	0.22	0.36	0.29	0.29	0.22	0.36
BLEU	0.89**	0.64*	0.88**	0.84**	0.92**	0.93**	0.85**	0.71*	0.83*	0.71*	0.83*	0.71*	0.21	0.21	0.28	0.21	-0.01	0.14	-0.08	0.21	-0.01	0.14	-0.08	0.21
METEOR	0.93**	0.64*	0.88**	0.84**	0.94**	0.79**	0.91**	0.86**	0.87**	0.71*	0.89**	0.71*	0.49	0.50	0.54	0.50	0.31	0.36	0.24	0.29	0.31	0.36	0.24	0.29
CHRF	0.95**	0.64*	0.90**	0.84**	0.96**	0.93**	0.91**	0.71*	0.88**	0.71*	0.89**	0.71*	0.48	0.50	0.52	0.50	0.31	0.29	0.23	0.36	0.31	0.29	0.23	0.36
CIDEe	0.75*	0.50	0.83**	0.69*	0.85**	0.79**	0.82*	0.71*	0.82*	0.57	0.83*	0.57	0.12	0.00	0.20	0.00	-0.03	0.07	-0.09	0.00	-0.03	0.07	-0.09	0.00
BERTScore	0.94**	0.71*	0.87**	0.62*	0.93**	0.71*	0.89**	0.93**	0.85**	0.79**	0.86**	0.79**	0.54	0.43	0.59	0.43	0.54	0.43	0.48	0.36	0.54	0.43	0.48	0.36
MOVERScore	0.97**	0.79**	0.93**	0.69*	0.97**	0.79**	0.93**	0.86**	0.87**	0.71*	0.88**	0.71*	0.55	0.50	0.60	0.50	0.46	0.43	0.39	0.36	0.46	0.43	0.39	0.36
ROUGE-we	0.95**	0.71*	0.94**	0.76**	0.98**	0.86**	0.95**	0.79**	0.90**	0.64*	0.91**	0.64*	0.50	0.50	0.55	0.50	0.45	0.43	0.38	0.36	0.45	0.43	0.38	0.36
EmbeddingAverage	0.79*	0.50	0.82*	0.69*	0.86**	0.79**	0.87**	0.71*	0.85**	0.57	0.86**	0.57	0.71*	0.57	0.75	0.57	0.56	0.50	0.51	0.43	0.56	0.50	0.51	0.43
VectorExtrema	0.80*	0.57	0.80*	0.76**	0.86**	0.86**	0.82*	0.64*	0.84**	0.64*	0.84**	0.64*	0.37	0.21	0.42	0.21	0.40	0.36	0.33	0.29	0.40	0.36	0.33	0.29
GreedyMatching	0.89**	0.64*	0.80*	0.69*	0.88**	0.79**	0.85**	0.71*	0.85**	0.71*	0.86**	0.71*	0.60	0.50	0.64	0.50	0.43	0.50	0.36	0.43	0.43	0.43	0.35	0.36
SummaQA	0.87**	0.57	0.85**	0.62*	0.91**	0.71*	0.93**	0.79**	0.87**	0.64*	0.89**	0.64*	0.24	0.21	0.30	0.21	0.43	0.43	0.35	0.36	0.43	0.43	0.35	0.36

Pearson's  $r$  and Kendall's  $\tau$  between intrinsic automatic metrics and extrinsic criteria. Significance is indicated by \* for p-values less than or equal to 0.05 and \*\* for p-values less than or equal to 0.01

## (2) Evaluating automatic metrics using extrinsic criteria:

- According to top-k system analysis, most automatic metrics fail to consistently and reliably quantify differences in usefulness between systems.



System-level Pearson correlations between intrinsic automatic metrics and proposed extrinsic metrics on top-k systems



# Conclusions

---

- An extrinsic evaluation framework to assess the usefulness of text summaries with a web-based platform to facilitate the data collection.
- A new human extrinsic evaluation dataset with 4k annotated articles.
- We find that summaries are generally useful in tasks that require a comprehensive understanding of an article. We also explore the connection between the usefulness and intrinsic properties of summaries.
- We re-evaluate 14 automatic metrics and discover that most of them fail to reflect the extrinsic metrics in classification and similarity tasks.